

AI WENT LIVE. THE TRAINING DIDN'T.

3 in 4

Nearly 3 in 4 organisations have deployed or are piloting AI

18%

Only 18% say most of their workforce got AI upskilling in the past year

+107%

Critical thinking enrolments. Complex problem-solving +126%; agentic workflows (an AI skill) +546%

Source: Coursera + Udemy, Global Skills Report 2026. The 3-in-4 and 18% figures cite Aon, 2026. Growth is year on year.

Access to Copilot or ChatGPT is not capability. Training completion counts who sat the course, not who now works differently. The real test is whether people still understand the process, question the evidence, verify root causes and own the decision when AI is in the loop. These five measures look at that.

1 Verification rate

- Of the AI outputs that fed a decision this month, how many were checked against the source data or record?
- Sample ten AI-assisted documents and trace each claim back.

Trap: Counting licences and logins as readiness.

OUTPUT

The share of AI-assisted decisions with the evidence checked, by team.

2 Challenge rate

- How often do people correct, reject or ask for more on an AI suggestion?
- Never challenging is a warning, not a win.

Trap: Reading 100% acceptance as proof the tool works.

OUTPUT

Logged challenges, and what they caught.

3 Root cause quality

- In AI-assisted investigations, is each root cause proven with evidence?
- Watch repeat deviations and “human error” closures, not closure speed.

Trap: Faster CAPA closure taken as better problem-solving.

OUTPUT

Verified root causes and repeat rate, before vs after AI.

4 Decision lead time

- Time from the first signal (alarm, trend, complaint) to a decision someone owns.
- AI should shorten the reading. People still make the call.

Trap: Timing how fast the report is generated, not the decision.

OUTPUT

A baseline and trend for two or three recurring decisions.

5 Named ownership

- Every AI-supported decision has a named person who signs it, and a written rule for when to override.
- Critical in GMP: the record shows who decided, not which tool.

Trap: “The system recommended it” in a batch record or investigation.

OUTPUT

An owner and override rule for each AI use case.



COMPLETION COUNTS PEOPLE. CAPABILITY CHANGES RESULTS.

Most AI programmes report the left column. Ask for the right one.

WHAT USUALLY GETS REPORTED	WHAT TO ASK FOR INSTEAD
Licences issued, monthly active users	Share of AI-assisted decisions with the evidence checked
Training hours, course completion	Who now works differently: one changed task per person, observed
Prompts written, documents generated	Challenges logged, and the errors they caught
Investigations closed faster	Root causes proven with evidence; repeat deviations
Reports produced in minutes	Decision lead time, from signal to an owned decision
“AI-enabled” processes	Processes that were stable and baselined before AI touched them

OPTIMISE BEFORE YOU AUTOMATE

AI on an unstable process learns the instability. Before an AI use case goes live, check the process is walked, standardised and baselined. Otherwise you can't tell whether AI helped.

TRAIN THINKING BEFORE TRUST

Teach people to question the answer: where did this come from, what would prove it wrong, what does the record say? Learners are already moving this way: critical thinking enrolments grew 107% in the report.

WARNING SIGNS: ADOPTION WITHOUT CAPABILITY

- ☐ Readiness is reported as licences, logins or hours
- ☐ Nobody can name a decision AI changed last month
- ☐ AI suggestions are almost never challenged
- ☐ Investigations close faster, but the same deviations return
- ☐ Operators first saw the tool at go-live
- ☐ No one owns the override when AI is wrong

Two or more ticks? Run the scorecard on page 3 with the team that uses the tool.

START THIS WEEK

- ☐ Pick one AI use case already live on your site.
- ☐ Pull five decisions it supported last month. For each, ask: was the evidence checked, was the output challenged, who signed it?
- ☐ Bring the answers to your next tier meeting. That is your readiness baseline.



AI READINESS SCORECARD ON ONE PAGE

Fill it in with the people who use the tool, for one AI use case. Score each measure from 1 (not measured) to 5 (measured, moving the right way, and acted on).

AI USE CASE

TEAM AND PROCESS IT SUPPORTS

MEASURE	HOW WE MEASURE IT	TODAY	TARGET	OWNER	SCORE
1 · Verification rate					1 2 3 4 5
2 · Challenge rate					1 2 3 4 5
3 · Root cause quality					1 2 3 4 5
4 · Decision lead time					1 2 3 4 5
5 · Named ownership					1 2 3 4 5

FIVE DECISIONS AI SUPPORTED LAST MONTH

TRAINING THAT WOULD MOVE THE LOWEST SCORE

DECISION	CHECKED?	CHALLENGED?	SIGNED	SKILL OR PRACTICE	WHO	BY WHEN

READING YOUR TOTAL (5 TO 25)

5 to 11	Adoption, not capability. Start with ownership and verification before adding use cases.
12 to 19	Partly ready. Fix the lowest score first, and train the thinking behind it.
20 to 25	Capability. Share how you measure it with the next site.